

# Introduction to Unsupervised Learning

Gabriel Maher

[gdmaher@stanford.edu](mailto:gdmaher@stanford.edu)

Alexander Ioannidis

[ioannidis@stanford.edu](mailto:ioannidis@stanford.edu)



Institute for Computational and Mathematical Engineering,  
Stanford University

# Unsupervised Learning

- *Unsupervised learning*: a set of statistical tools for data for which only features/inputs are available
  - We have  $X$ 's but no associated labels  $Y$
  - Goal: discover interesting patterns/properties of the data
    - e.g. for visualizing or interpreting high-dimensional data

# Unsupervised Learning

- Sample applications:
  - Given a collection of text documents, identify sets of documents about the same topic
  - Given high-dimensional facial images, find a compact representation as inputs for a facial recognition classifier



(AT&T Laboratories  
Cambridge)

# Unsupervised Learning

- Why is unsupervised learning challenging?
  - Exploratory data analysis – goal is not as clearly defined
  - Difficult to assess performance – “right answer” unknown
  - Working with high-dimensional data

# Unsupervised Learning



- Two approaches:
  - *Cluster analysis*
    - For identifying homogenous subgroups of samples
  - *Dimensionality Reduction*
    - For finding a low-dimensional representation to characterize and visualize the data

# Cluster Analysis & K-means

# Clustering

- *Clustering*: a set of methods for finding subgroups within the data set
  - Observations should share common characteristics within the same subgroup, but differ across subgroups
  - Groupings are determined from the data itself – differs from classification



# Clustering



Data sets from: <http://cs.joensuu.fi/sipu/datasets/>



# Clustering



Data sets from: <http://cs.joensuu.fi/sipu/datasets/>


# Clustering

- Types of clustering models
  - Centroid-based clusters
  - Hierarchical clustering
  - Model-based clustering
    - Each cluster is represented by a parametric distribution
    - Data set is a mixture of distributions
  - Hard vs. soft/fuzzy clustering
    - Hard: Observations divided into distinct clusters
    - Soft: Observations may belong to more than one cluster

# K-means Clustering

- Groups data into K distinct clusters
  - Each of K clusters is defined by a centroid vector
    - Centroid vector: mean of all the samples assigned to cluster
  - Each observation assigned to a single cluster (nearest centroid)
  - Requires number of clusters K as input
  - “Good clustering” minimizes within-cluster variation
    - “Similarity” measured by Euclidean distance

# K-means Clustering



\*Some of the figures in this presentation are taken from *"An Introduction to Statistical Learning, with applications in R"* (Summerer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani

Figure 10.5 , ISL 2013\*

# K-means Clustering




Figure 10.5 , ISL 2013

# K-means Clustering

- Centroids minimize within-cluster variation

$$J = \frac{1}{n} \sum_{i=1}^K \sum_{x \in c_i} |x - \mu_i|^2$$

- Centroids (cluster centers):  $\mu_i = \frac{1}{|c_i|} \sum_{x \in c_i} x$


- Minimization is a combinatorial optimization problem
  - Solve for local minimum using an iterative method

# K-means Algorithm\*

- 1) Select initial set of centroids
- 2) Partition data by assigning each sample to cluster associated with its nearest centroid
- 3) Compute new centroids within each cluster
- 4) Repeat 2 and 3 until convergence
  - “converged” when centroids stabilize and samples do not move between clusters

\* also known as “Lloyd’s algorithm”


# K-means Algorithm



Pick initial centroids




# K-means Algorithm




Pick initial centroids  
Assign initial clusters

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters  
Update centroids

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters  
Update centroids

# K-means Algorithm



Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters

# K-means Algorithm




Pick initial centroids  
Assign initial clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters  
Update centroids  
Reassign clusters  
Converged



# K-means Algorithm

- Poor initialization leads to poor clustering




# Initializing Centroids

- Random selection of  $K$  samples
- Random partition of data
- Select  $K$  points that are mutually “far apart”
- Domain-knowledge
  - Identify samples that we expect to belong to different clusters
- Initialize using results from another clustering method


# How many clusters?

- K-means requires we provide  $K$  (# clusters) as input
  - We may have domain-knowledge to inform choice of  $K$
  - Otherwise, choice of  $K$  is determined from data


$K = 2, J = 110.935$




$K = 3, J = 30.5285$




$K = 4, J = 4.0934$



$K = 5, J = 2.6754$



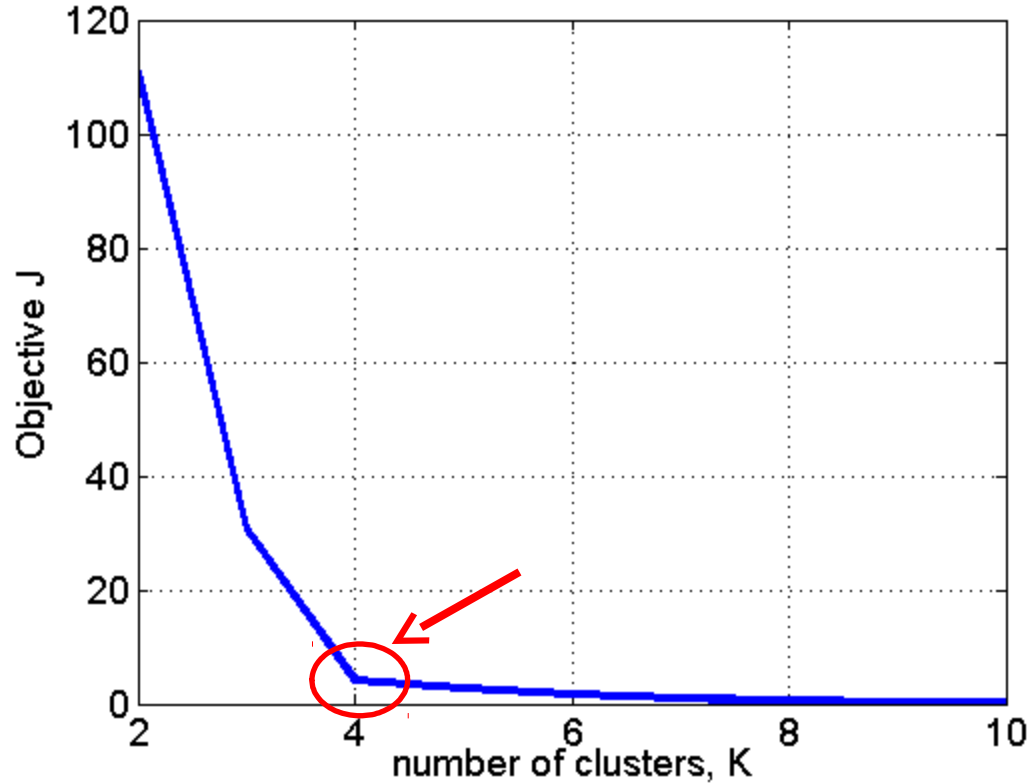
$K = 6, J = 1.7164$



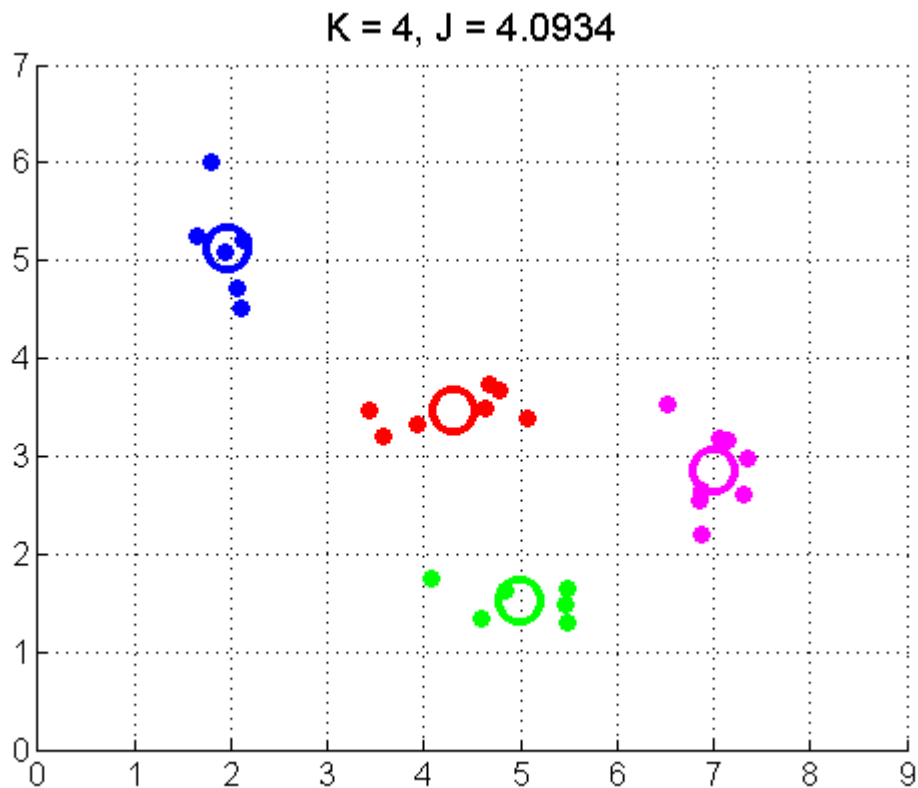
# How many clusters?

- Can't just pick value of  $K$  that minimizes objective  $J$ 
  - $J$  decreases monotonically with increasing  $K$
- Heuristic method:
  - For each candidate value  $K$ ,
    - Compute  $K$ -means clustering  $M$  times, find minimum objective  $J_K$
  - Find “elbow” in objective curve ( $K$  vs  $J_K$ )

# How many clusters?



# How many clusters?




# K-means algorithm

- Advantages
  - Easy to implement
  - Often converges in a small number of iterations
  - Can be applied on data with a large number of features
- Disadvantages
  - K is input parameter
  - Iterative algorithm returns local minimum\*




# K-means algorithm



# K-means algorithm

- Advantages
  - Easy to implement
  - Often converges in a small number of iterations
  - Can be applied on data with a large number of features
- Disadvantages
  - K is input parameter
  - Iterative algorithm returns local minimum\*
  - Assumes all clusters spherical and approximately the same size\*


# K-means algorithm



# K-means algorithm

- Advantages
  - Easy to implement
  - Often converges in a small number of iterations
  - Can be applied on data with a large number of features
- Disadvantages
  - K is input parameter
  - Iterative algorithm returns local minimum\*
  - Assumes all clusters spherical and approximately the same size\*
  - Sensitive to outliers\*

# K-means algorithm



# K-means algorithm

- Advantages
  - Easy to implement
  - Often converges in a small number of iterations
  - Can be applied on data with a large number of features
- Disadvantages
  - K is input parameter
  - Iterative algorithm returns local minimum\*
  - Assumes all clusters spherical and approximately the same size\*
  - Sensitive to outliers\*
  - **\*some disadvantages handled by variants of K-means**

# Example:

## Image Segmentation/Compression




# Image Segmentation/Compression

- Image  $\rightarrow$  pixels  $\rightarrow$  RGB vectors (colors)
- Apply K means to collection of RGB vectors
  - One RGB vector per pixel
  - Clusters represent similar colors
- Replace each pixel with its associated centroid
  - Results in image with only K different colors




# Image Segmentation/Compression

K = 3




# Image Segmentation/Compression

K = 4




# Image Segmentation/Compression

K = 5




# Image Segmentation/Compression

K = 20



# Image Segmentation/Compression

$K = 20$



Questions?



# Introduction to Supervised Learning

Gabriel Maher  
[gdmaher@stanford.edu](mailto:gdmaher@stanford.edu)

Institute for Computational and Mathematical Engineering,  
Stanford University



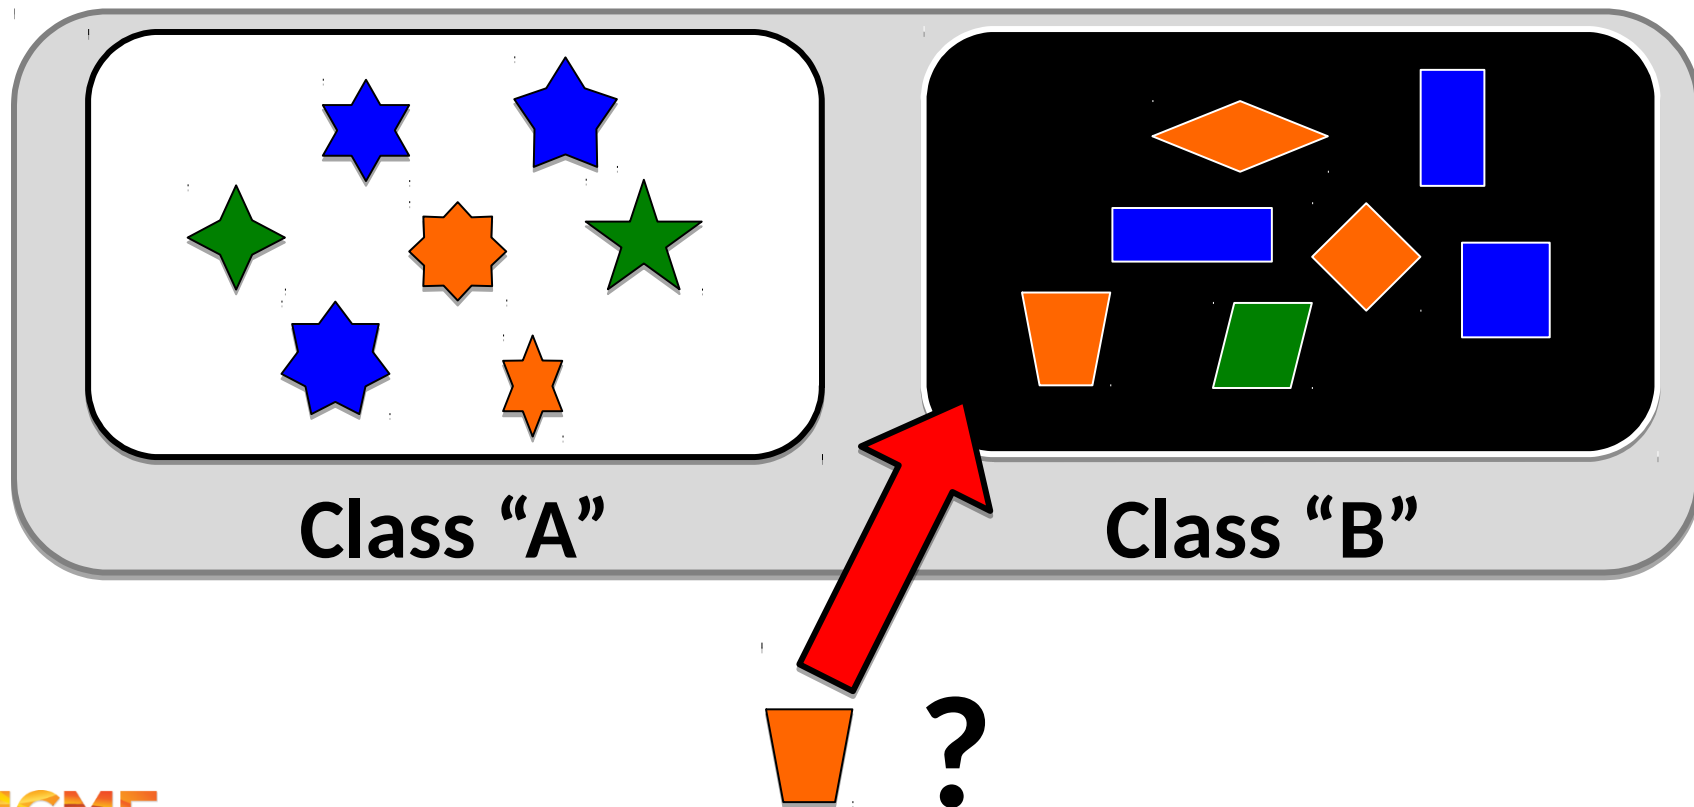
# Supervised Learning

# Supervised Learning


- Recall framework  $Y = f(X) + \epsilon$
- Supervised learning methods:
  - “Learn by example”
  - Build a model  $\hat{f}$  using set of labeled observations

$$\left(X^{(1)}, Y^{(1)}\right), \dots, \left(X^{(n)}, Y^{(n)}\right)$$

# Training Data



# Training data



\*Some of the figures in this presentation are taken from *"An Introduction to Statistical Learning, with applications in R"* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani

Figure 2.2 , ISL 2013

# Supervised Learning

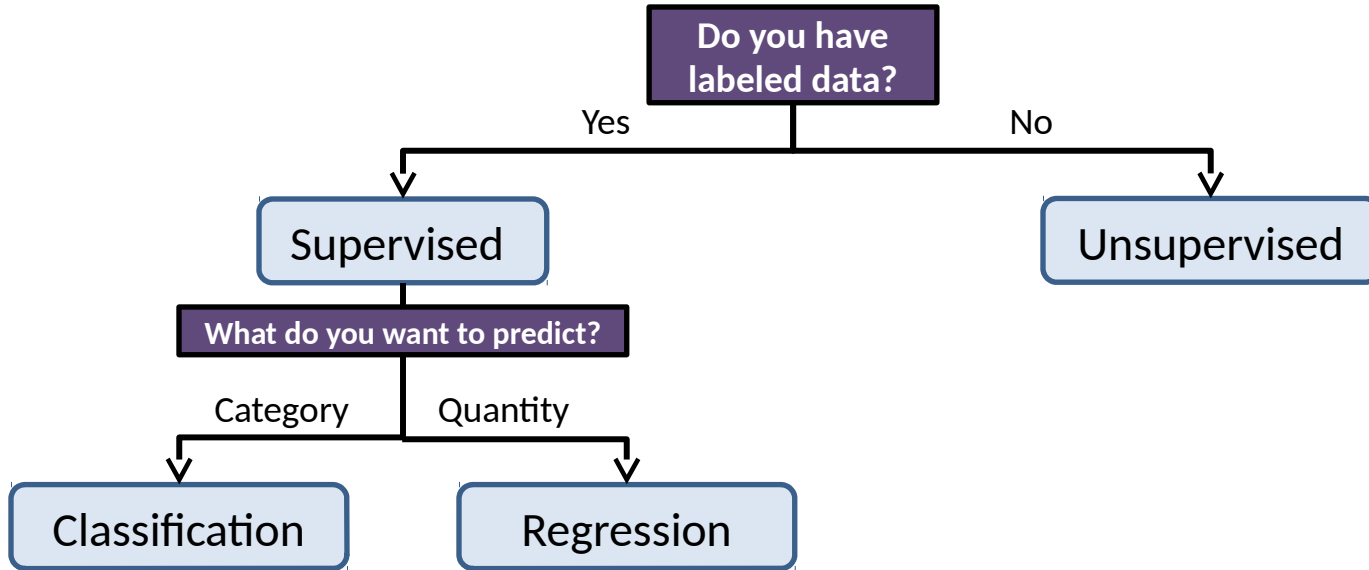
- Supervised learning algorithms
  - Pick “best” estimate  $\hat{f}$  from among family of functions
- Example: Linear regression
  - Select best fit to *training data* within set of all linear functions

$$f(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_d X_d$$

# Classification and Regression

- Supervised learning problems belong to two types
  - Regression: the output  $Y$  is quantitative
  - Classification: the output  $Y$  is qualitative/categorical

# Types of Algorithms



# Assessing Model Accuracy



# Measuring Performance: Regression

- Loss function: quantifies cost of errors
- e.g. Mean squared error (MSE)
  - Common measure of prediction accuracy in regression

$$MSE = \frac{1}{n} \sum_{i=1}^n \left( \hat{y}^{(i)} - y^{(i)} \right)^2$$

- Penalizes large errors more than small errors



# Measuring Performance: Regression

- Our goal: build model that *generalizes*
  - We want to minimize errors on unseen “test data,” not on the training data
  - e.g. predicting *future* stock prices vs. past stock prices
- Want to minimize *expected* loss
  - Problem: Can’t just minimize loss on training data

# Challenge: Overfitting

- *Overfitting*: learning the random variation in the data rather than the underlying trend
- Characteristic of overfitting:
  - Good performance on previously-seen (training) data, but poor performance on new data

# Challenge: Overfitting



Figures 2.4 and 2.6 , ISL 2013

# Measuring Performance

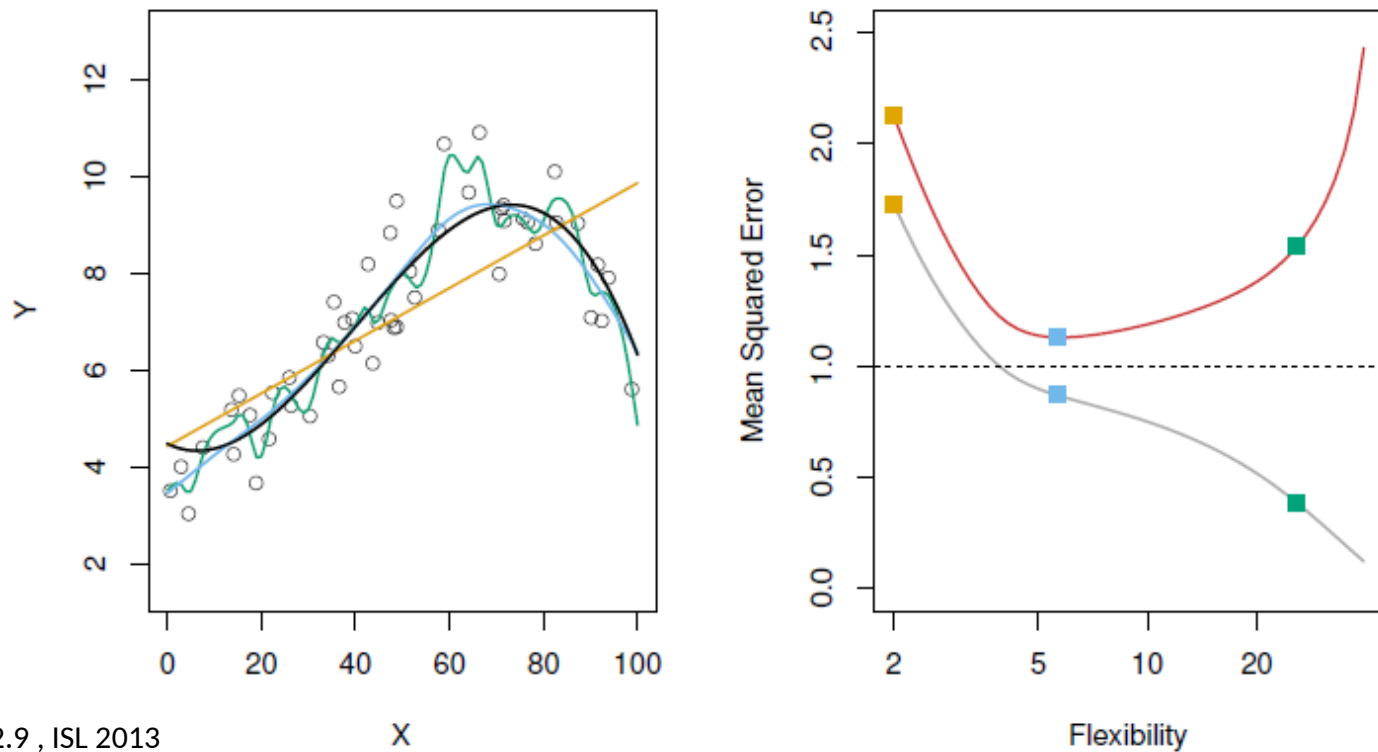


Figure 2.9 , ISL 2013

# Measuring Performance

- How do we estimate test error to find a good model?
- *Cross-validation:*  
a set of techniques for using the training data to estimate generalization error

# Data sets

- *Training data*
  - Set of observations used to learn the model
- *Validation data*
  - Set of observations used to estimate error for parameter-tuning or model selection
- *Test data*
  - Set of observations used to measure performance on unseen data
  - These data are **not** available to the algorithm during learning process

# Trade-off: Bias vs. Variance




- U-shaped test error due to two competing properties of learning methods:

$$\mathbb{E} [\text{test error}] = \text{var}(\hat{f}) + \text{bias}(\hat{f})^2 + \text{var}(\epsilon)$$




- $\text{var}(\hat{f})$ : *variance* of estimator
- $\text{bias}(\hat{f})$ : *bias* of estimator



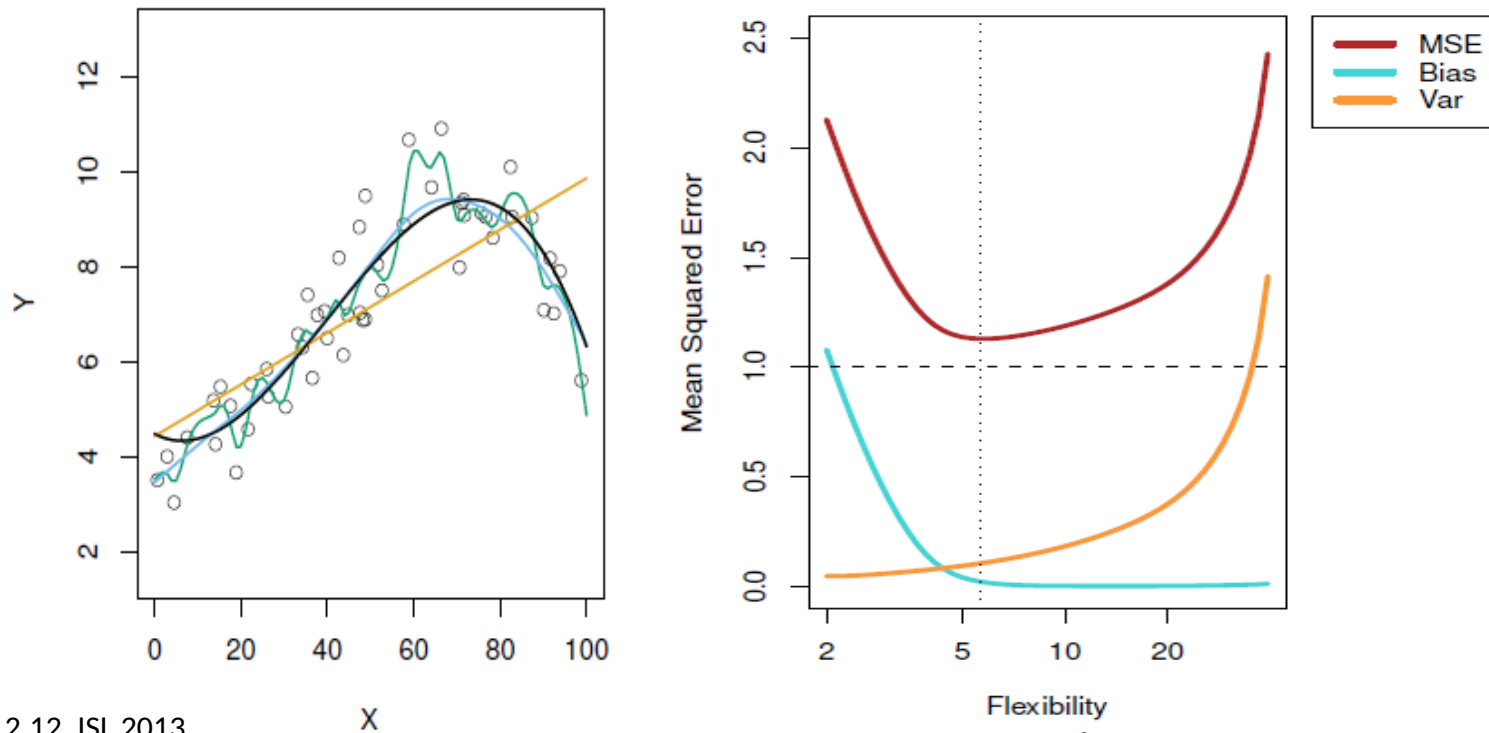
# Trade-off: Bias vs. Variance



# Trade-off: Bias vs. Variance



# Trade-off: Bias vs. Variance



Figures 2.9, 2.12, ISL 2013

# Trade-off: Bias vs. Variance

- Flexible method
  - Can achieve closer fit to underlying system (low bias),
  - But risks learning model that is too dependent on training data (high variance)
- Simpler method
  - May not accurately model underlying system (high bias),
  - But less dependent on training data (low variance)
- Trade-off
  - Easy to achieve low variance/high bias *or* high variance/low bias,
  - But hard to achieve low variance *and* bias

# Regression: Linear Regression

# Linear Regression

- *Linear Regression*: a simple supervised learning method, used to predict quantitative output values
  - Many machine learning methods are generalizations of linear regression
  - Example to illustrate key concepts in supervised learning

# Simple Linear Regression

- Output  $Y$  and single predictor  $X$  with linear relationship between  $X$  and  $Y$ :

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- Model parameters:
  - $\beta_0$  intercept
  - $\beta_1$  slope

# Simple Linear Regression




Figure 3.1 , ISL 2013



# Simple Linear Regression

- $\beta_0$  and  $\beta_1$  are unknown  $\rightarrow$  estimate their values using training data

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$$

- Pick  $\hat{\beta}_0, \hat{\beta}_1$  such that the model is a “good fit” to training data

$$Y^{(i)} \approx \hat{\beta}_0 + \hat{\beta}_1 X^{(i)}, \quad i = 1, \dots, n$$

# Simple Linear Regression

- How do we estimate the coefficients (“fit the model”)?
- What makes model a “good fit” to the data?

# Least Squares

- Typically, fit of the model to the data is measured using *least squares*
- Mean squared error:

$$MSE = \frac{1}{n} \sum_{i=1}^n \left( Y^{(i)} - \hat{Y}^{(i)} \right)^2$$

# Least Squares

- “Fit model” - coefficients  $\hat{\beta}_0, \hat{\beta}_1$  minimize MSE:

$$\min_{(\hat{\beta}_0, \hat{\beta}_1)} \left[ \frac{1}{n} \sum_{i=1}^n \left( Y^{(i)} - \left( \hat{\beta}_0 + \hat{\beta}_1 X^{(i)} \right) \right)^2 \right]$$

- Explicit formulas for solution:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X^{(i)} - \bar{X})(Y^{(i)} - \bar{Y})}{\sum_{i=1}^n (X^{(i)} - \bar{X})^2} \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

# Simple Linear Regression




Figure 3.1 , ISL 2013

# Multiple Linear Regression




Figure 3.4 , ISL 2013

# Multiple Linear Regression

- Multiple linear regression: more than one predictor variable used to predict response

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_d X_d + \epsilon$$

# Least Squares

- Solve for coefficient estimates (“fit model”) using least squares fit on training data:

$$\hat{\beta} = \arg \min_{\beta} \|Y - X^T \beta\|^2 \quad X = \begin{bmatrix} 1 & X^{(1)T} \\ & \dots \\ 1 & X^{(n)T} \end{bmatrix} \quad \hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \dots \\ \hat{\beta}_d \end{bmatrix}$$

- Use normal equations to solve for  $\hat{\beta}$ :  $Y = \begin{bmatrix} Y^{(1)} \\ \dots \\ Y^{(n)} \end{bmatrix}$

$$X^T X \hat{\beta} = X^T Y \quad \rightarrow \quad \hat{\beta} = (X^T X)^{-1} X^T Y$$



# Linear Regression

- Advantages:
  - Simple model
  - Interpretable coefficients
  - Can obtain good results with small data sets
  - Many variations/extensions
- Disadvantages:
  - Model may be too simple to make accurate predictions over large range of values
  - Poor extrapolation
  - Sensitive to outliers in data – least squares error

Questions?

# Classification: Logistic Regression

# Classification

- Regression – predicting quantitative response  $Y$ 
  - For some applications, response variable is qualitative or categorical
- Classification: predicting a qualitative response
  - Assign each observation to a class / category
  - e.g. K-nearest neighbor classifier from Lecture 1

# Classification and Regression

- Classification and regression are closely related
- Classification as regression:
  - Predict the probability an observation belongs to each class, assign to class with highest probability

# Logistic Regression

- *Binary classification*:  $Y$  takes on two values (“0” or “1”) corresponding to two classes
- Logistic regression models binary classification as

$$Pr(Y \text{ belongs to class } 1 \mid X)$$

- Threshold to obtain class decisions
- Modified Linear regression for probabilities in  $[0, 1]$

# Logistic Regression

- Logistic (sigmoid) function bounds output  $Y \in [0, 1]$
- Logistic function  $\sigma(z) = \frac{e^z}{1 + e^z} \left( = \frac{1}{1 + e^{-z}} \right)$ 
  - S-shaped curve
  - Always takes values in (0, 1) -> valid probabilities
- Logistic Regression model

$$Pr(Y = 1 \mid X) = \sigma(\beta_0 + \beta_1 X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

# Logistic Regression

- Model parameters  $\beta_0$  and  $\beta_1$  must be estimated using the training data
  - For linear regression, use *least-squares*
- Fit model parameters using different cost function
  - Binary cross-entropy



# Logistic Regression

- *Binary cross-entropy*

$$L = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - 1) \log(1 - \hat{y}^{(i)}) - y^{(i)} \log \hat{y}^{(i)}$$

- *Bad prediction:*  $y = 1, \hat{y} = 0 \rightarrow L = \infty$   
 $y = 0, \hat{y} = 1 \rightarrow L = \infty$
- *Good prediction:*  $y = 1, \hat{y} = 1 \rightarrow L = 0$   
 $y = 0, \hat{y} = 0 \rightarrow L = 0$

# Multiple Logistic Regression

- We can extend logistic regression to the case of multiple predictor variables:

$$\begin{aligned} Pr(Y = 1 \mid X) &= \sigma(\beta_0 + \beta_1 X_1 + \cdots + \beta_d X_d) \\ &= \frac{e^{(\beta_0 + \beta_1 X_1 + \cdots + \beta_d X_d)}}{1 + e^{(\beta_0 + \beta_1 X_1 + \cdots + \beta_d X_d)}} \end{aligned}$$

# Logistic Regression

- Advantages:
  - Extension of linear regression
  - No hyperparameters to tune
- Disadvantages:
  - Can not model complex decision boundaries
  - Must be formulated as binary classification problem

# Summary

- *Supervised learning* – learning from examples
- *Linear regression* – simple, interpretable model for predicting quantitative response
- *Logistic regression* – regression method used to predict probabilities for binary classification
  - *Maximum likelihood method*: technique for estimating parameter values

Questions?